

Aryan Nagwekar

FOR COTIVITI · AGENTIC AI RESEARCHER

M.S. Electrical & Computer Engineering, Rutgers University · aryan.nagwekar@gmail.com · (732) 543-9282 · [linkedin.com/in/aryan-nagwekar](https://www.linkedin.com/in/aryan-nagwekar) · aryan-junction.vercel.app

► See it running (90-sec demo): youtu.be/eGIZYUU7_A

Healthcare runs on **trust and auditability** — so agentic AI can only be deployed here if every action it takes is **governed, risk-scored, and explainable**. That's exactly what I research and build: not just capable AI agents, but the systems that make them **safe, auditable, and trustworthy** enough to act in a regulated domain. It's the precise intersection of **agentic AI** and **healthcare-grade accountability** this role lives at.

WHAT I'VE BUILT — AGENTIC AI, WITH AUDITABILITY AT THE CORE

AegisAI — a runtime firewall for AI agents. Intercepts an agent's shell, file, API, and MCP tool calls *before they execute*; risk-scores each action; blocks secret exfiltration and destructive commands; routes risky actions to human approval; scans MCP tool metadata for poisoning; and exports audit logs for review. → governance and auditability for autonomous AI — the prerequisite for trusting agents in a compliance-heavy domain like healthcare.

AgentLab — observability & control for agentic systems. A runtime to watch, govern, and *replay* what AI agents do — every action attributed to an agent, every failure explained in plain English, the whole run reconstructable like a recording. ~600 automated tests; event-sourced architecture. → research-grade visibility into agent behavior, and a deterministic audit trail of every decision.

Agentic, tool-using data pipeline. Built a pipeline that ingests and de-duplicates multi-source data, generates embeddings for retrieval, classifies with a *hybrid rule-based + LLM* approach, and exports auditable SQL/CSV for human verification. → applied agentic AI — retrieval, LLM reasoning, and a human-checkable audit trail, end to end.

WHY I FIT THE AGENTIC AI RESEARCHER ROLE

Agentic AI systems	My whole body of work is autonomous agents and the runtime that governs them — not a side interest.
LLMs, retrieval & evaluation	Embeddings/retrieval + hybrid rule-based/LLM inference, with deterministic checks over the outputs.
Auditability & risk control (healthcare-critical)	Every system I build risk-scores actions, keeps a human in the loop, and exports audit logs by design.
Research rigor & self-direction	Took on an open problem — governing autonomous agents — and shipped a working, tested runtime solo.
Communicating to technical & non-technical audiences	A core design goal of my work is making agent behavior legible in plain English, not stack traces.

HOW I THINK ABOUT AGENTIC AI IN HEALTHCARE

- In a regulated domain, an agent's *auditability* matters as much as its capability — you have to be able to prove what it did and why.
- The safest pattern isn't "human reviews everything" or "agent does everything" — it's *risk-scored autonomy*: act freely on the safe 95%, escalate the risky 5%.
- Trust in an agentic system starts with evaluation and explainability — a system that can't explain itself can't be deployed where money and care are on the line.

I'd love to bring this to Cotiviti's agentic-AI research work. · Aryan Nagwekar · aryan.nagwekar@gmail.com · (732) 543-9282 · aryan-junction.vercel.app